

Integration of the camera depth image

3EY4 - Lab 10 - Bonus

AV#23

Danial Noori Zadeh

400367734

Hassan Al Masmoum

400330500

Kevin Le

400385350

April 9, 2024

Table of Contents

| | |
|---|------------------|
| <i>Problem Identification</i> | <i>3</i> |
| <i>Solution Overview</i> | <i>3</i> |
| <i>Design steps</i> | <i>4</i> |
| Step 1 – launching the camera node & listening to the camera messages. | 4 |
| Step 2 – collecting camera information data | 5 |
| Step 3 – receive and de-projecting camera depth data. | 6 |
| Verification of deprojection algorithm in MATLAB..... | 6 |
| Step 4 – Integration into LiDAR callback function | 8 |
| <i>Testing & Verification</i> | <i>10</i> |
| Static test | 10 |
| AEV in action..... | 10 |

Problem Identification


We were tasked with integrating the depth cloud data published by the camera with the LiDAR to detect obstacles that are not on the plane of the LiDAR. Implied in this problem is the requirement of efficiency and rate matching between the two sources of information (LiDAR and depth camera). A mismatch of sampling time or the resolution of can cause mis-operation of the self-driving algorithm.

Solution Overview

The camera publishes two topics that are key to our solution. One is the camera info that publishes information needed to convert depth (d) and pixel location (u, v) to xyz coordinates and later to r, θ polar coordinate to refine the LiDAR published ranges. The second important message topic is `/camera/depth/image_rect_raw` that contains the actual depth readings. We use CVBridge library to convert the image raw data to a matrix, but we were unable to pip install pyrealsense2 python wrapper to extract the xyz data, so, we manually computed them from the camera intrinsic data.

The solution uses a range of interest (ROI) over both u and v axis of the image pixels to reduce the unnecessary calculations. Another method improve efficiency is to filter the extracted xyz points based on a x and y window to avoid reading ground points and detection of objects that are out of the trajectory of the car. The results are confirmed by publishing the refined scan ranges and the performance was significant improved by optimizing the ROI (from 6 Hz publish rate to 12 Hz).

Figure 4-1. Depth Field of View to Depth Map illustration



Note:

1. As the scene's distance from the depth module increases, the invalid depth band decreases in the overall depth image. Overall depth image is invalid depth band plus valid depth map.

Design steps

Step 1 – launching the camera node & listening to the camera messages.

We used the `rs_camera.launch` file provided by the camera to `realsense2` package, enabled the IR for extra precision and reduced the publish rate to 15 fps. However, it is recommended by the realsense documentation that both the depth and RGB camera operate at the same frequency and as high as possible, and if needed apply “decimation” filter to bring down the rate. We have the launch file is called in the self-driving package under `experiment.launch` as shown below.

```

_av23@av23-desktop:~$ roslaunch realsense2 camera
demo_pointcloud.launch          rs_camera.launch              rs_multiple_devices.launch
demo_t265.launch               rs_d400_and_t265.launch       rs_rgbd.launch
opensource_tracking.launch      rs_d435_camera_with_model.launch rs_rtabmap.launch
rs_aligned_depth.launch         rs_from_file.launch           rs_t265.launch
_av23@av23-desktop:~$ roslaunch realsense2_camera

```

```

77  <!-- Launch the camera node -->
78  <include file="$(find realsense2_camera)/launch/rs_camera.launch">
79    <arg name="depth_width"      default="848"/>
80    <arg name="depth_height"     default="480"/>
81    <arg name="enable_depth"     default="true"/>
82    <arg name="depth_fps"        value = "15"/>
83
84  </include>

```

The important camera depth data are published the highlighted topics, which we set up in the initialization of the navigation.py file as shown below. We need to import **Image** and **CameraInfo** datatypes from `sensors.msg` to launch these publishers and subscribers. The “`refined_lidar_ranges`” is the new topic published to debug the corrected ranges from the program.

```

102  rospy.Subscriber(depth_image_topic, Image, self.imageDepthCallback, queue_size=1)
103  rospy.Subscriber(depth_info_topic, CameraInfo, self.imageDepthInfoCallback, queue_size=1)
104  # Publisher for refined data
105  self.refined_lidar_pub = rospy.Publisher("refined_lidar_ranges", LaserScan, queue_size=5)
106

```

```

_av23@av23-desktop:~$ rostopic list
/camera/color/camera_info
/camera/color/image_raw
/camera/color/metadata
/camera/depth/camera_info
/camera/depth/image_rect_raw
/camera/depth/metadata

```

Step 2 – collecting camera information data

Looking into a sample message from the camera_info topic we can see the following fields which contain many useful information such intrinsic matrix (K) of the camera.

```
av23@av23-desktop:~$ rostopic echo -n 1 /camera/depth/camera_info
header:
  seq: 6381
  stamp:
    secs: 1711601916
    nsecs: 53215504
  frame_id: "camera_depth_optical_frame"
height: 480
width: 848
distortion_model: "plumb_bob"
D: [0.0, 0.0, 0.0, 0.0, 0.0]
K: [421.70062255859375, 0.0, 425.8759765625, 0.0, 421.70062255859375, 239.4052734375, 0.0, 0.0, 1.0]
R: [1.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 1.0]
P: [421.70062255859375, 0.0, 425.8759765625, 0.0, 0.0, 421.70062255859375, 239.4052734375, 0.0, 0.0, 0.0, 1.0, 0.0]
binning_x: 0
binning_y: 0
roi:
  x_offset: 0
  y_offset: 0
  height: 0
  width: 0
  do_rectify: False
```

we then extracted the important information from the K field and stored them globally for the depth projection in the imageDepthCallback function.

```
643 def imageDepthInfoCallback(self, cameraInfo):
644     try:
645         self.intrinsics = True
646         self.img_width = cameraInfo.width
647         self.img_height = cameraInfo.height
648         self.ppx = cameraInfo.K[2]
649         self.ppy = cameraInfo.K[5]
650         self.fx = cameraInfo.K[0]
651         self.fy = cameraInfo.K[4]
652         #print("CALIBRATION: found camera Info\n")
653         #print(self.fx, ' ', self.fy, ' ', self.ppx, ' ', self.ppy, '\n')
654
655     except CvBridgeError as e:
656         print(e)
657         return
658
```

Pinhole Camera Model

- **Intrinsic** camera parameters:

$$K \triangleq \begin{bmatrix} \frac{f}{s_x} & 0 & c_x \\ 0 & \frac{f}{s_y} & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

- **Extrinsic** camera parameters:

$$T_w^c \triangleq [R_w^c \mid t_w^c]$$

Step 3 – receive and de-projecting camera depth data.

The xyz coordinates of the points in the camera depth image can be computed from the depth and pixel location (u and v). We used a reduced ROI to improve the efficiency of the deprojection process. We filtered out the ground points by filter the points based on their y-value. These settings can be changed in the params.yaml file.

$$Z \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

$$u = \frac{f_x}{z} x + c_x \Rightarrow x = \frac{z}{f_x} [u - c_x]$$

$$v = \frac{f_y}{z} y + c_y \Rightarrow y = \frac{z}{f_y} [v - c_y]$$

$$z = z$$

```

199 # Camera Parameters
200 use_camera: 1 # use the camera depth data (1) else 0 -> do not
201 min_depth: 100 # mm
202 max_depth: 5000 # mm
203 cam_baselink_offset: 270 #mm
204
205 # the following are defined in number of pixels
206 image_width: 848
207 image_height: 480
208
209 roi_u_lower: 250
210 roi_u_upper: 548
211 roi_v_lower: 150
212 roi_v_upper: 470
213
214 y_cam_max: 100 #mm
215 y_cam_min: -150 #mm
216 x_cam_max: 600 #mm
217 x_cam_min: -600 #mm

```

```

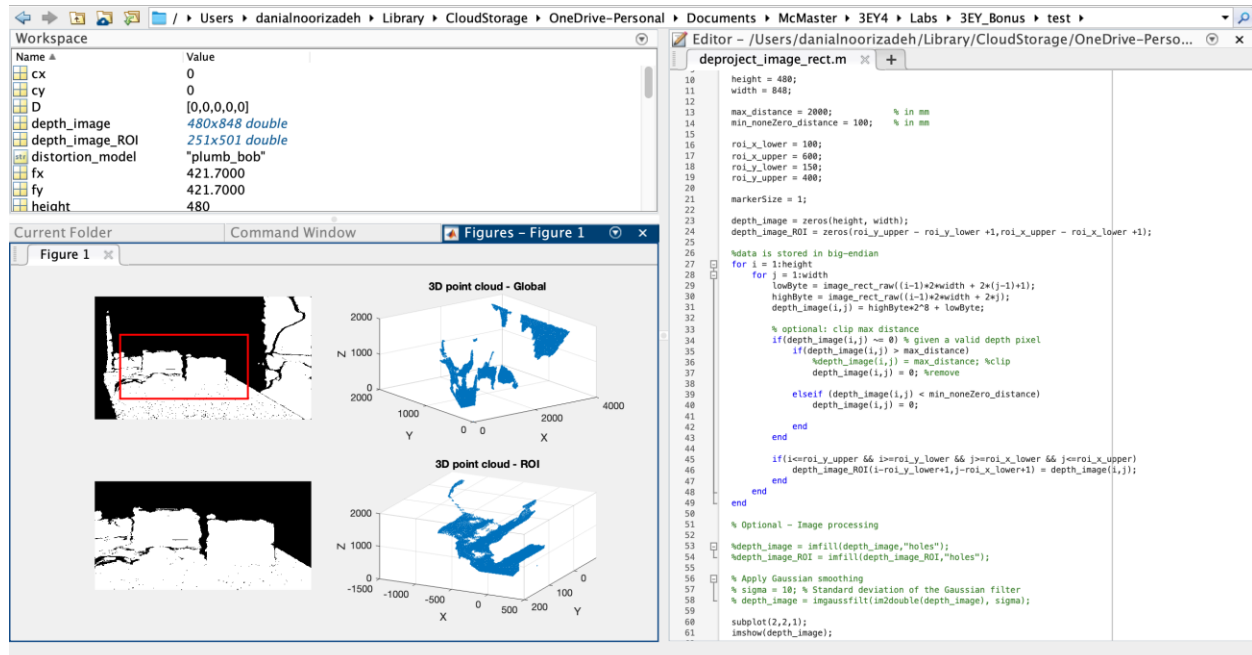
507 # start of camera callback functions
508 def imageDepthCallback(self, data):
509
510     try:
511         cv_image = self.bridge.imgmsg_to_cv2(data, encoding)
512         # height, width = cv_image.shape
513         if self.intrinsics:
514             # Convert all pixels to points within the region of interest (ROI)
515             v_range = np.arange(self.roi_v_lower, self.roi_v_upper + 1)
516             u_range = np.arange(self.roi_u_lower, self.roi_u_upper + 1)
517             u_grid, v_grid = np.meshgrid(u_range, v_range)
518
519             depths = cv_image[v_grid, u_grid]
520             depths[depths > self.max_depth] = 0
521             depths[depths < self.min_depth] = 0
522             depths = depths.astype(float)
523
524             # deproject depth to point
525             x = (depths / self.fx) * (u_grid - self.ppx)
526             y = (depths / self.fy) * (v_grid - self.ppy)
527             z = depths
528
529             # filter the points out of the range
530             mask = (y < self.y_cam_max) & (y > self.y_cam_min) & (z > 0) & (z < self.max_depth) & (x < self.x_cam_max) & (x > self.x_cam_min)
531             x = x[mask]
532             y = y[mask] # we don't need the y component anymore
533             z = z[mask]
534             z = z + self.cam_baselink_offset
535
536             x = x.astype(float)
537             y = y.astype(float)
538             z = z.astype(float)
539
540             # Calculate the polar coordinate in the camera frame from the point in the xz plane
541             self.cam_ranges = np.sqrt((x**2) + (z**2)) # Euclidean distance
542             self.cam_angles = np.arctan2(z, x) # Angle in radians
543
544             #OPTIONAL: Store xyz coordinate
545             #self.depth_points = np.stack([x,y,z], axis=1)
546
547     except CvBridgeError as e:
548         except Exception as e:

```

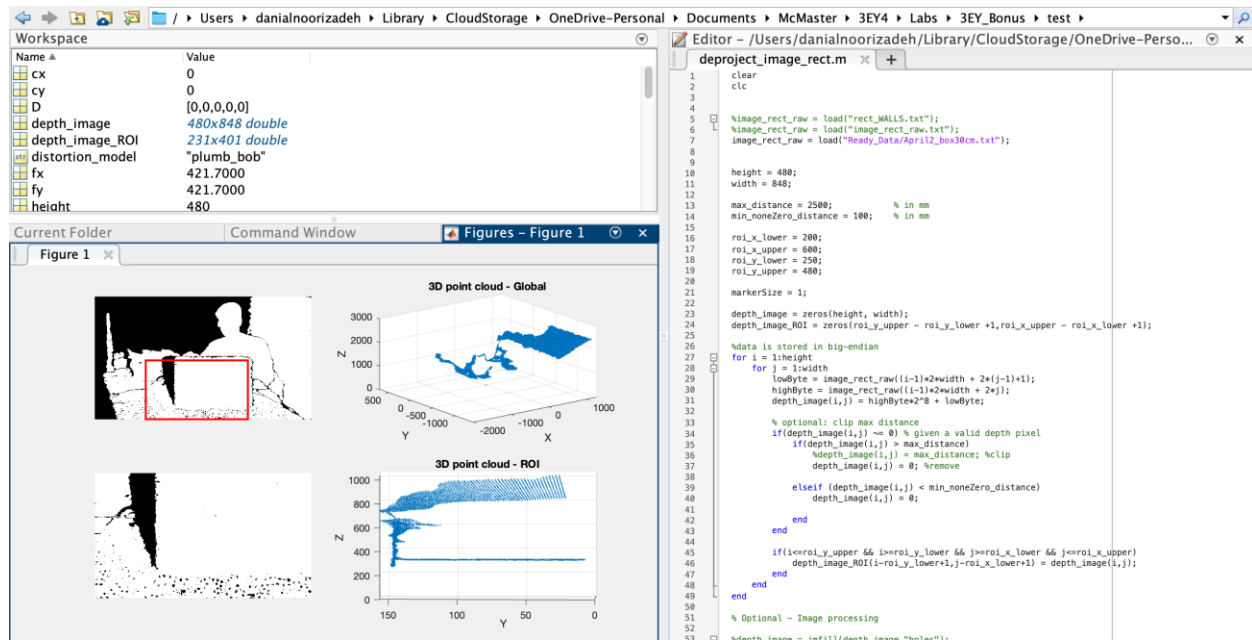
The z value is compensated by self.cam_baselink_offset to translate the points into the LiDAR frame and then the points are expressed in polar coordinates. an important step before conversion to polar using numpy square root and inverse tangent is that the arrays have to be set as float type otherwise the function will exhibit erroneous behavior.

Verification of deprojection algorithm in MATLAB

To verify the correctness of the deprojection function we made a MATLAB script that converts the image message to an image matrix, iterates over the depth points within the ROI and then computes the xyz coordinates. We used experimental data collected by placing a block that is out of the sight of the LiDAR at different distances away from the AEV. We monitored the output point cloud and tuned the ROI and y limits.



The red square displays the ROI setting. The advantage of performing the test in MATLAB is that we can use image processing tools such as hole filling filters to improve the readings.



Step 4 – Integration into LiDAR callback function

The LiDAR data needs to be refined before any pre-processing or navigation algorithm is run on it. We would only integrate the camera info if the global `use_camera` variable is set to true, the camera intrinsic data is collected and the `cam_ranges` have been extracted. We then pass the LiDAR ranges, camera ranges and corresponding angles to the `refine_lidar` method and create a new `LaserScan` message for the refined ranges.

```

327 def lidar_callback(self, data):
328
329     ranges = data.ranges
330
331     t = rospy.Time.from_sec(time.time())
332     self.current_time = t.to_sec()
333     dt = self.current_time - self.prev_time
334
335     self.prev_time = self.current_time
336     sec_len = int(self.heading_beam_angle/data.angle_increment)
337
338     if (self.use_camera and self.intrinsics and (self.cam_ranges is not None) and (self.cam_angles is not None)):
339         # refine the ranges
340         ranges = self.refine_lidar(ranges, self.cam_ranges, self.cam_angles)
341
342         # Create a new LaserScan message for the refined ranges
343         refined_ranges_msg = LaserScan()
344         refined_ranges_msg.header.stamp = rospy.Time.now()
345         refined_ranges_msg.header.frame_id = data.header.frame_id # Assuming the frame ID is the same as the original LiDAR data
346         refined_ranges_msg.angle_min = data.angle_min
347         refined_ranges_msg.angle_max = data.angle_max
348         refined_ranges_msg.angle_increment = data.angle_increment
349         refined_ranges_msg.time_increment = data.time_increment
350         refined_ranges_msg.scan_time = data.scan_time
351         refined_ranges_msg.range_min = data.range_min
352         refined_ranges_msg.range_max = data.range_max
353         refined_ranges_msg.ranges = ranges
354
355         # Publish the refined ranges
356         self.refined_lidar_pub.publish(refined_ranges_msg)
357
358
359     proc_ranges, mod_ranges = self.preprocess_lidar(ranges)
360
361
362 > if self.drive_state == "normal": ~
493
494 > elif self.drive_state == "backup": ~
531
532 > else: ~
584
585     # Publish to driver topic
586     drive_msg = AckermannDriveStamped()

```

We first cast the ranges from LiDAR to a list, so it becomes mutable. We then find the closest lidar index that the `cam_angles` correspond to by dividing the angle by the angle increment of the LiDAR. It's important we add $\pi/2$ to the angles to compensate for the LiDAR x axis being 90 degrees rotated with respect to the camera frame. We have already compensated for the translation between the two frames by adding the offset to the z-values.

The program will then iterate over the ranges in the camera, access the corresponding angle in the LiDAR frame and only overwrite the ranges if the camera suggests a smaller value. The `corrected_point` is a debugging variable used to store the number of corrections made to the LiDAR ranges.


```
190 # Refine LiDAR Data
191 def refine_lidar(self, ranges_LiDAR, cam_ranges, cam_angles):
192
193     # Convert ranges_LiDAR tuple to a list
194     refined_ranges = list(ranges_LiDAR)
195
196     # Convert camera angles to LiDAR indices
197     lidar_indices = np.round((cam_angles + np.pi/2.0)/ self.ls_ang_inc).astype(int)
198
199     corrected_points = 0
200     # Iterate through each LiDAR index and compare the ranges
201     for i in range(len(lidar_indices)):
202         # Check if the camera range is smaller than the LiDAR range at the corresponding index
203         if float(cam_ranges[i]/1000.0) < refined_ranges[lidar_indices[i]]:
204             # Update the LiDAR range with the camera range
205             refined_ranges[lidar_indices[i]] = float(cam_ranges[i]/1000.0)
206             corrected_points = corrected_points + 1
207
208     #print("corrected ", corrected_points, " times\n")
209     # OPTIONAL: Convert the refined_ranges list back to a tuple if needed
210     #refined_ranges = tuple(refined_ranges)
211
212     return refined_ranges
213
```

Testing & Verification

Static test

The published refined LiDAR was compared to the original LiDAR in rviz. We further verify the corrected points are coming from the camera's depth cloud by enabling the depth cloud topic of the camera in rviz.

AEV in action

The vehicle is able to avoid stationary objects assuming they are already in the camera's line of sight; however, due to the difference in timing between LiDAR and the refined camera info the vehicle cannot avoid the objects that it did not have enough time to recognize.